

Business Intelligence with Apache Druid

Technical Tutorial

Data Science Conference Europe
15-19th November 2020

About mentor

- Vladimir Ivković
- Teaching Assistant at
Faculty of Technical Sciences,
University of Novi Sad
- Computing and Control Department
- Chair for Applied Computer Science
- Group for Data Science and
Information Systems



Prerequisites

- Software and tools
 - Docker & Docker Compose
 - cURL
 - JDK 8 (optionally)
 - DB client (optionally)
- Knowledge
 - SQL
 - Database Design

Tutorial Agenda

Introduction to Data
Warehouse Systems

Exercise 1: Running
first Druid cluster and
simple file ingestion

Q&A Session

10 mins

20 mins

~30 mins

~30 mins

~10 mins

Apache Druid
Concepts and
Capabilities

Exercise 2: Data
ingestion from source
database and querying

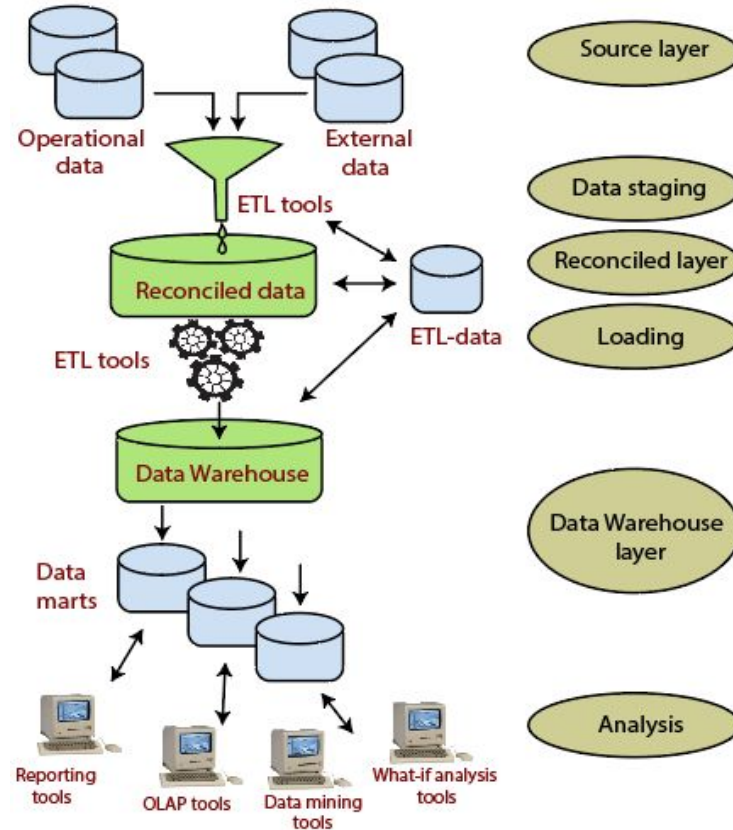
Data Warehouse - Introduction

- **Data Warehousing** is process for collecting and managing data from varied sources to provide meaningful business insights
- **Data Warehouse** is a central repository of information that can be analyzed to make more informed decisions
- Heterogeneous data sources
 - transactional systems
 - relational databases
 - all sorts of files and documents
- DW supports analytical reporting, structured and/or ad hoc queries, and decision making

Data Warehouse - Main Features

- Subject Oriented
 - DW is designed around “subjects” rather than processes.
 - DW offers information regarding a theme instead of companies' ongoing operations. These subjects can be sales, marketing, distributions, etc.
- Integrated
 - the establishment of a common unit of measure for all similar data from the various sources
- Nonvolatile
 - once entered into the data warehouse, data should not change
- Time Variant
 - data warehouse's focus on change over time
 - Historical data for trend analysis

Data Warehouse - Architecture



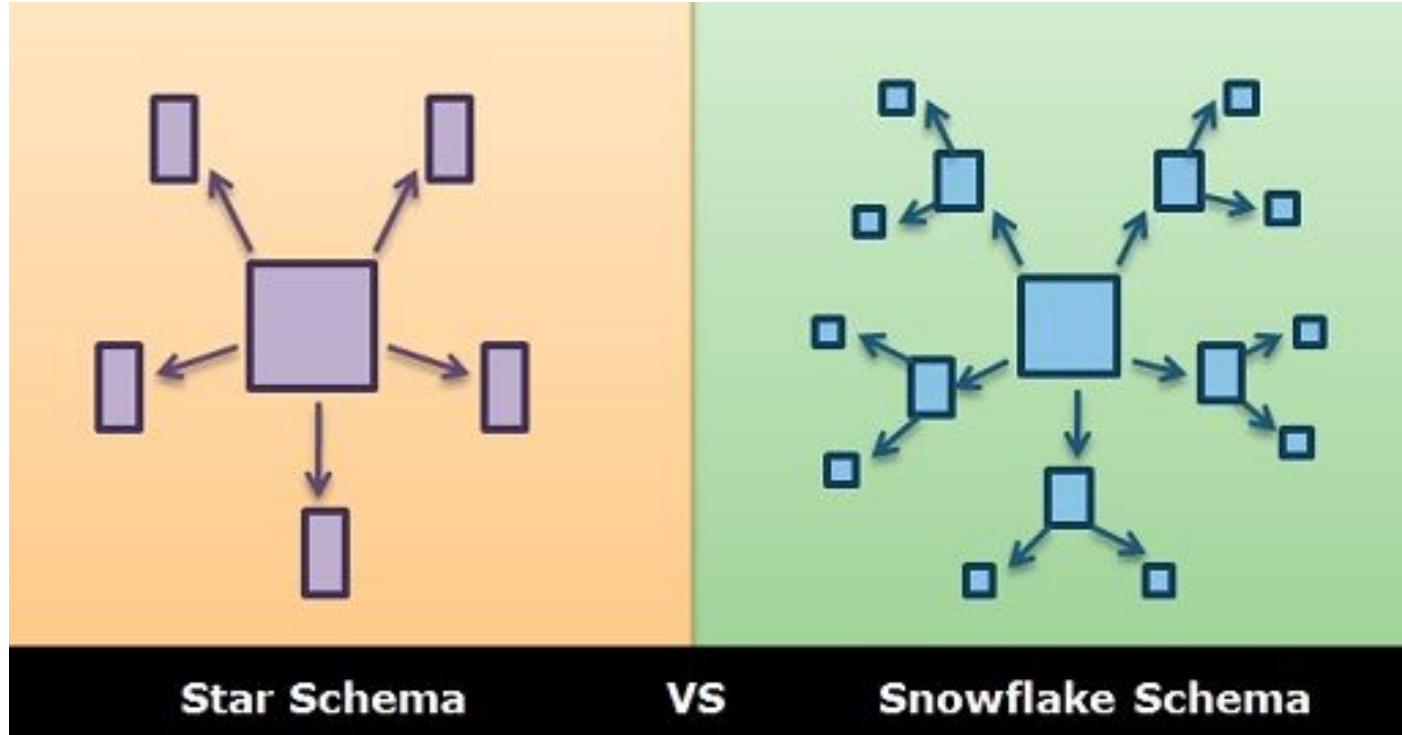
Three-Tier Architecture for a data warehouse system

source: <https://www.javatpoint.com/data-warehouse-architecture>

Data Warehouse Schema Design

- OLTP vs. OLAP
- Bulk insert - no update - frequent and complex queries
- Data warehouses often use denormalized or partially denormalized schemas to optimize query and analytical performance.
- Dimensional modeling approach
 - Dimensions - provide the “who, what, where, when, why, and how” context surrounding a business process event; descriptive context
 - Facts - the measurements that result from a business process event and are almost always numeric

Data Warehouse Schema Design



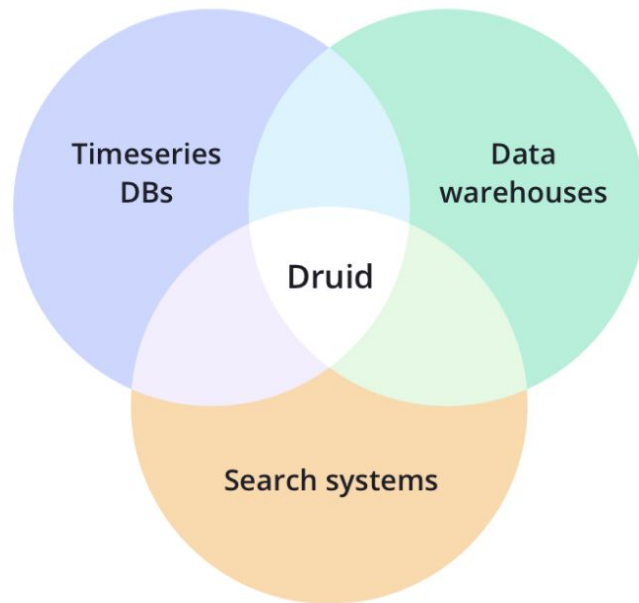
source: <https://techdifferences.com/difference-between-star-and-snowflake-schema.html>

Traditional Data Warehouse Systems

- Relational databases - primarily data source
- Structured data
- Relational database - choice for DW database
- Methodology matured in 90s
 - The Data Warehouse Toolkit by Ralph Kimball (John Wiley and Sons, 1996)
 - Building the Data Warehouse by William Inmon (John Wiley and Sons, 1996)
- Indexing - optimizing query execution
- Materialized views
 - used to precompute and store aggregated data such as the sum of sales
 - eliminates the overhead associated with expensive joins and aggregations for a large or important class of queries

Apache Druid - Introduction

- *Open source distributed data store*
- Combines ideas from data warehouses, time-series databases, and search systems
- High performance real-time analytics database for a broad range of use cases



Apache Druid - Key Features

- Column-oriented storage
- Native search indexes
- Streaming and batch ingest
- Flexible schemas
- Time-optimized partitioning
- SQL support
- Horizontal scalability
- Easy operation
- Initial requirements
 - Arbitrary queries
 - Scalability: trillions events/day
 - Interactive: low latency queries
 - Real-time: data freshness
 - High availability
 - Rolling upgrades
- Initial motivation
 - Business intelligence queries
 - Arbitrary slicing and dicing of data
 - Interactive real-time visualizations on complex data streams

Apache Druid - Common Use Cases

- Clickstream analytics (web and mobile analytics)
- Risk/fraud analysis
- Network telemetry analytics (network performance monitoring)
- Server metrics storage
- Supply chain analytics (manufacturing metrics)
- Application performance metrics
- **Business intelligence / OLAP**

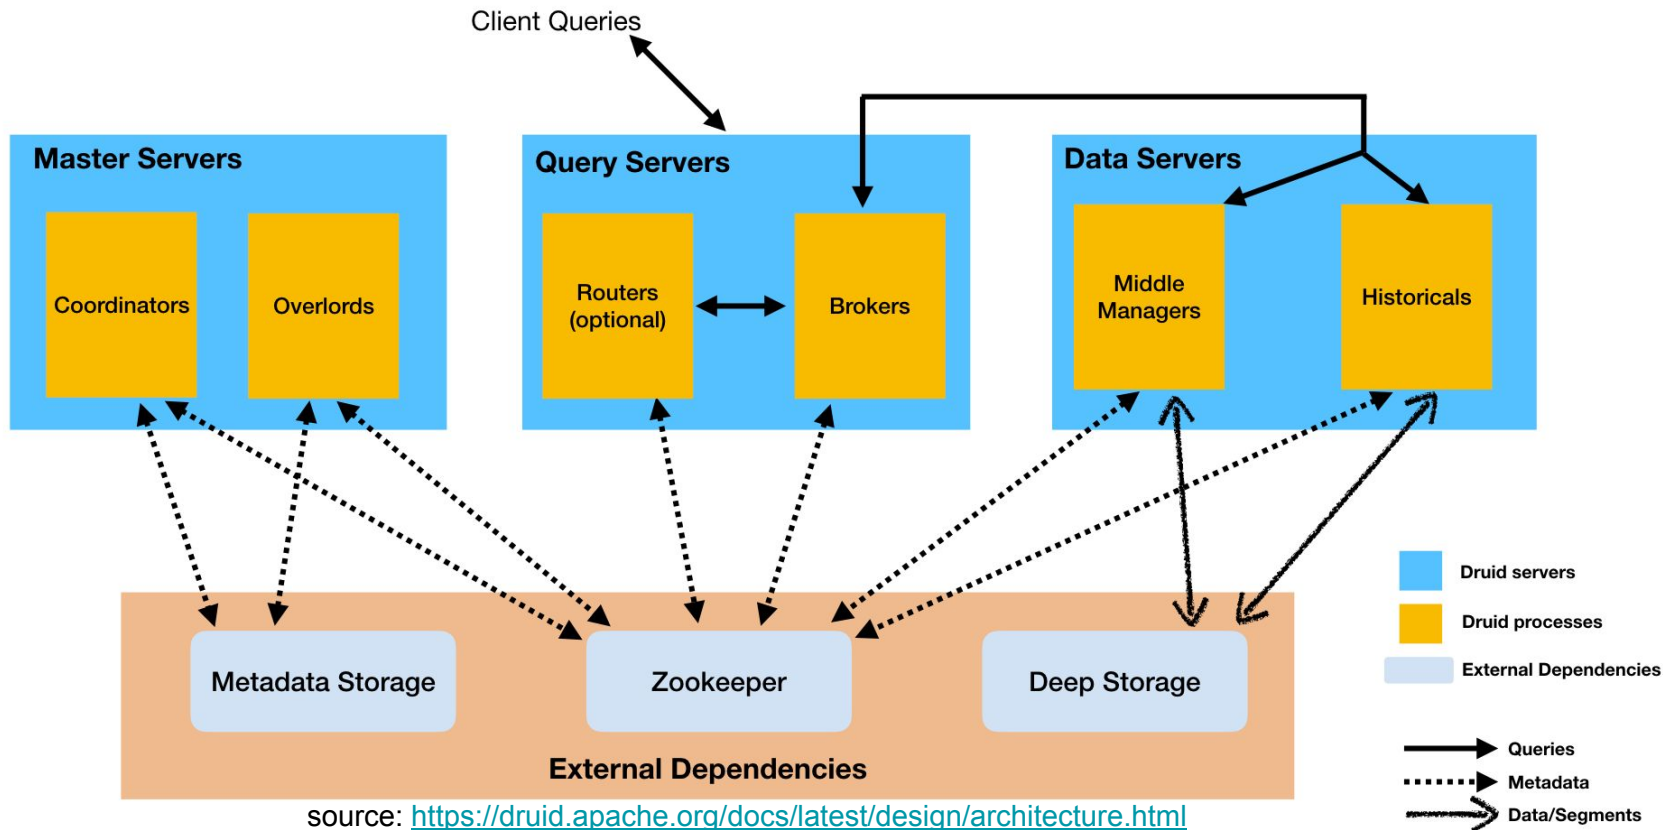
Apache Druid - Process Types

- *Druid architecture* - multi-process, distributed, cloud-friendly
- **Coordinator** - manage data availability on the cluster.
- **Overlord** - control the assignment of data ingestion workloads.
- **Broker** - handle queries from external clients.
- **Router** - are optional processes that can route requests to Brokers, Coordinators, and Overlords.
- **Historical** - store queryable data.
- **MiddleManager** - responsible for ingesting data.

Apache Druid - Server Types

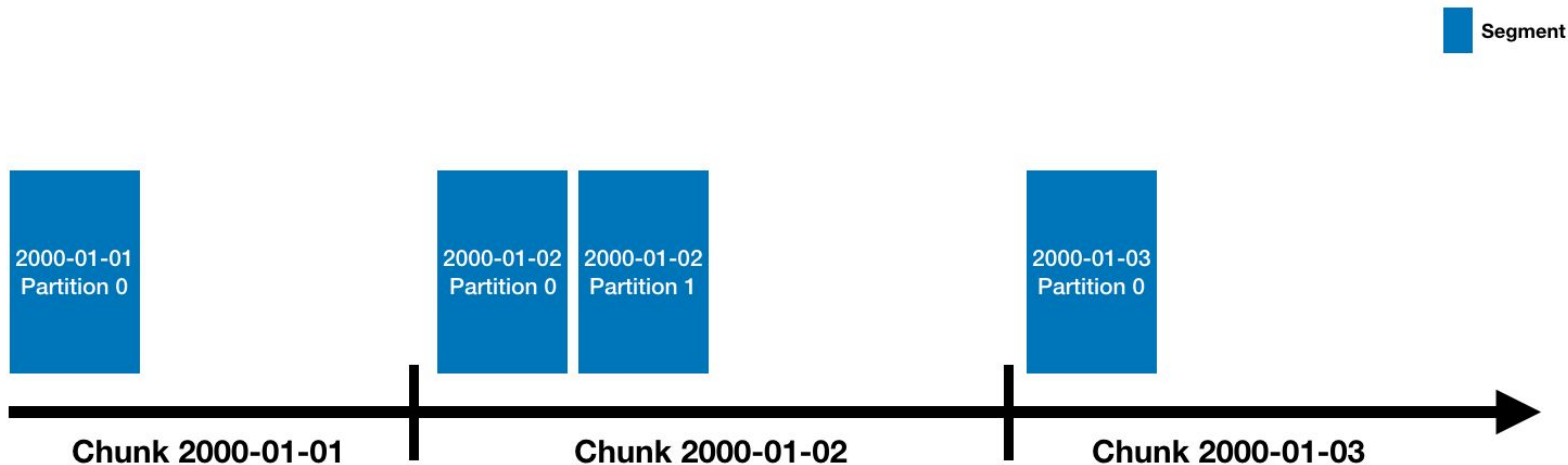
- It is suggested organizing processes into server types
- **Master**
 - runs *Coordinator* and *Overlord* processes
 - manages data availability and ingestion
- **Query**
 - runs *Broker* and optional *Router* processes
 - handles queries from external clients
- **Data**
 - runs *Historical* and *MiddleManager* processes
 - executes ingestion workloads and stores all queryable data

Apache Druid - Architecture



Apache Druid - Storage design

- Data is stored in ***datasources***
- Datasources are partitioned by time in ***chunks***
- Chunk contains one or more ***segments***
- Segment is a single file



Apache Druid - Ingestion

- Loading data in Druid is called **ingestion** or **indexing**
 - consists of reading data from a source system and creating segments based on that data
- Streaming ingestion methods - supervisor types
 - Kafka
 - Kinesis
 - Tranquility
- Batch ingestion methods - task types
 - Native batch - simple
 - Native batch - parallel
 - Hadoop

Apache Druid - Ingestion

- *Primary timestamp* - used for partitioning and sorting data
- *Dimensions* - columns stored as-is and can be used for any purpose
- *Metrics* - columns stored in an aggregated form
- *Rollup* - form of summarization or pre-aggregation used to minimize the amount of raw data that needs to be stored
- *Partitioning* - secondary partitioning using a particular dimension will improve locality
- **Ingestion specs** - JSON file consists of three main parts
 - **dataSchema** - definition of datasource, primary timestamp, dimensions, metrics, transformation and filtering
 - **ioConfig** - description of source system, connection method and data parsing
 - **tuningConfig** - controls various tuning parameters specific to each ingestion method

Apache Druid - Datasource

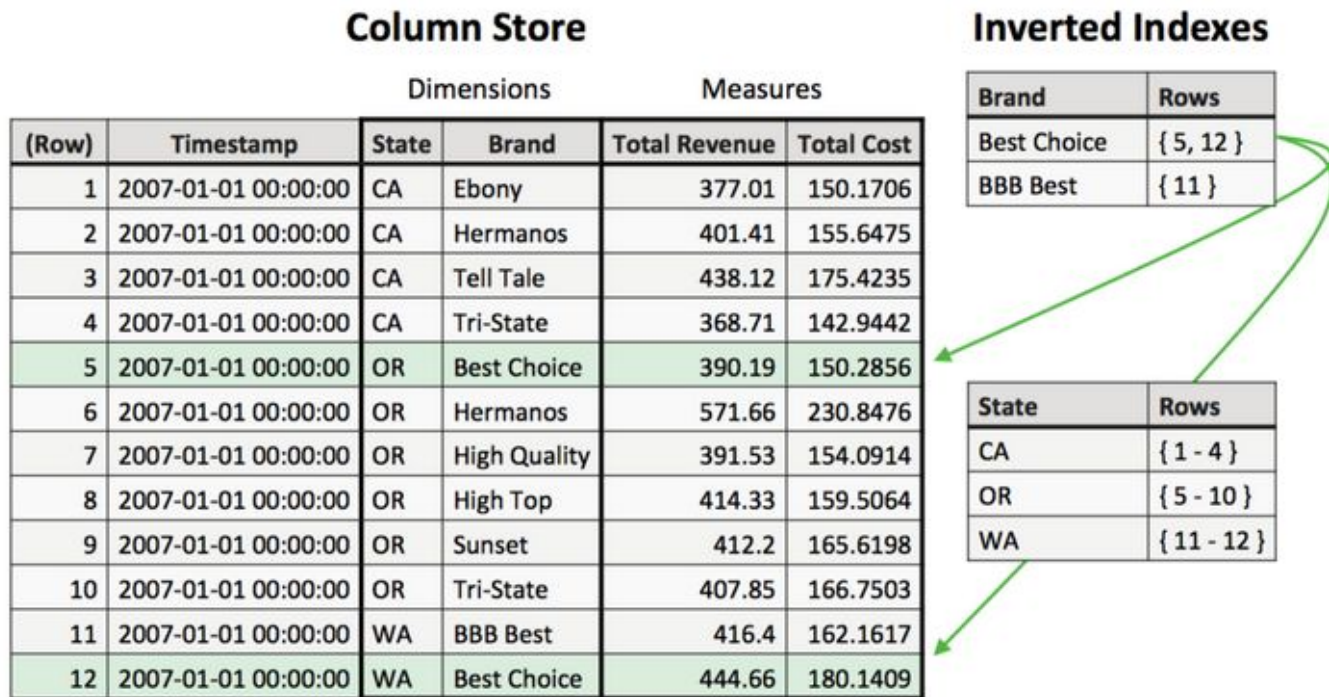
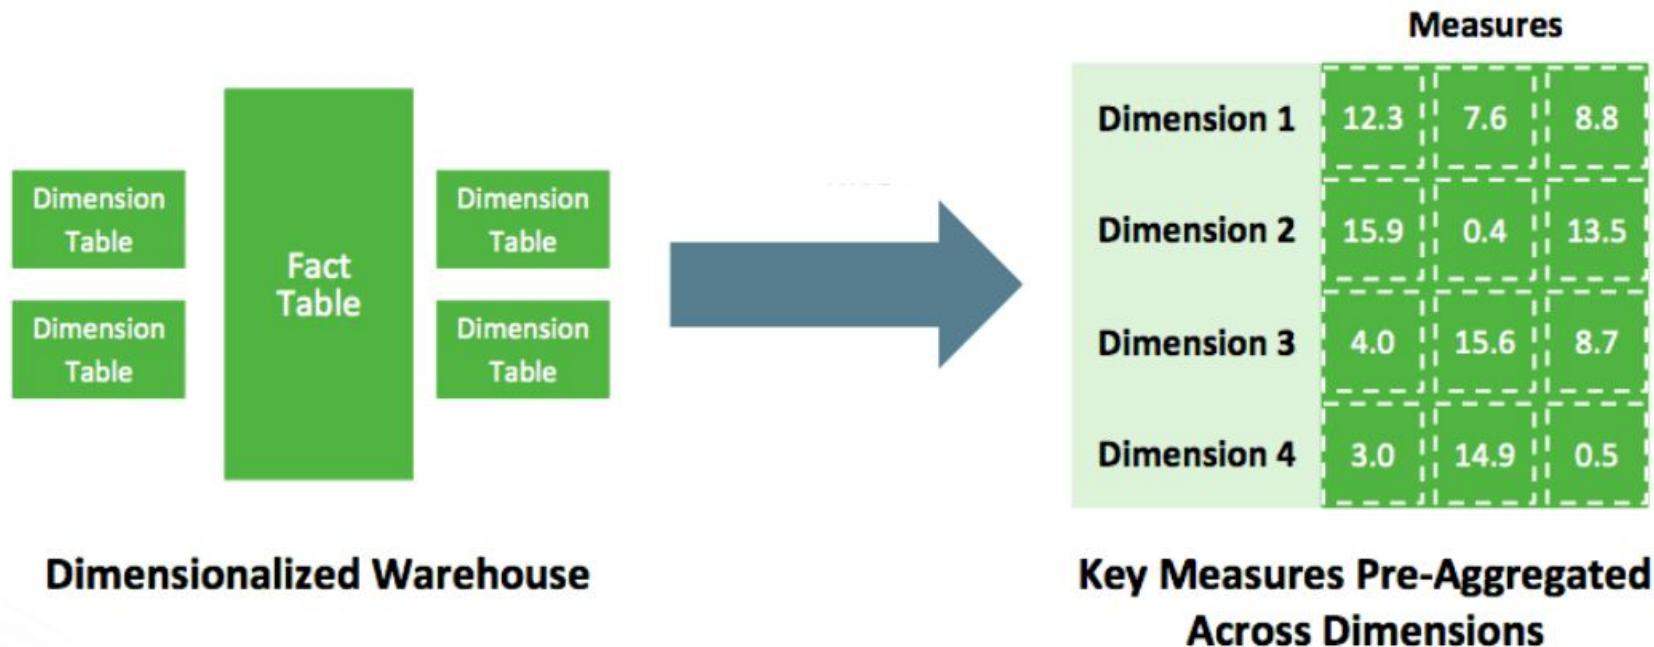


Figure 1: Druid combines the best qualities of a column store and inverted indexing

Dimensionalized DW vs Druid datastore



Apache Druid - Standalone Installation

- Prereqs
 - Linux, Mac OS X, or other Unix-like OS (Windows is not supported)
 - Java 8, Update 92 or later (8u92+)
- Install Druid
 - [Download](#) binaries
 - Extract Druid
- Start up Druid services
 - Single machine configuration - `./bin/start-micro-quickstart`
- Open the Druid console
 - Visit <http://localhost:8888>

Apache Druid - Data Loader (1/2)

- Click **Load data** from the Druid console header
- Select the **Local disk** tile and then click **Connect data**.
- Enter the following values:
 - **Base directory:** quickstart/tutorial/
 - **File filter:** wikticker-2015-09-12-sampled.json.gz
- Click **Apply**.
- Click **Next: Parse data**.
- With the JSON parser selected, click **Next: Parse time**.
- Click **Next: Transform**, **Next: Filter**, and then **Next: Configure schema**, skipping a few steps. Disable **Rollup**.

Apache Druid - Data Loader (2/2)

- Click **Next: Partition** to configure how the data will be split into segments. In this case, choose **DAY** as the **Segment granularity**.
- Click **Next: Tune** and **Next: Publish**.
- Let's change the default name from *wikiticker-2015-09-12-sampled* to *wikipedia*.
- Click **Next: Edit spec** to review the ingestion spec we've constructed with the data loader.
- Once you are satisfied with the spec, click **Submit**.
- Open **Ingestion** section.
- Once a task is completed, open **Datasources** section.
- Finally, go to **Query** section.

Apache Druid - Querying

- Apache Druid supports two query languages:
 - Druid SQL - powered by a parser and planner based on Apache Calcite
 - native queries - JSON-based query language
- Query examples for *wikipedia* datasource
 - `SELECT count(*) FROM wikipedia WHERE isAnonymous = TRUE`
 - `SELECT count(*) FROM wikipedia WHERE "comment" LIKE '%clean%'`
 - `SELECT page, cityName, countryName, isAnonymous
FROM "wikipedia" WHERE channel LIKE '%sr.%'`
 - `SELECT "countryName", count(*) "cnt" FROM "wikipedia"
GROUP BY "countryName" ORDER BY "cnt" DESC LIMIT 10`
 - `SELECT "user", count(*) "cnt" FROM "wikipedia"
WHERE "isRobot" = FALSE GROUP BY "user"
ORDER BY "cnt" DESC LIMIT 10`
- Druid console includes query editor with query building options

Apache Druid - Docker Compose

- Prereqs
 - Docker & Docker Compose
 - Download or clone <https://github.com/vladimirivkovic/druid-tutorial>
- Compose file - *docker-compose.yml*
 - a container for each Druid service
 - Zookeeper service
 - PostgreSQL container as the metadata store and OLTP database
 - Metabase as a graphical UI tool
- Configuration
 - Using environment variables, e.g. Druid extensions
 - Segments stored in a shared directory
- Launching the cluster
 - `docker-compose up`

Apache Druid - File Ingestion

- Sample ingestion specs for local file
 - On each container /quickstart/tutorial directory
- Several ways to submit ingestion task
 - Loading data with a spec (via console)
 - Load Data > Edit Spec > Submit
 - Loading data with a spec (via command line)
 - `bin/post-index-task --file quickstart/tutorial/wikipedia-index.json --url http://localhost:8081`
 - Loading data without the script
 - `curl -X 'POST' -H 'Content-Type:application/json' -d @quickstart/tutorial/wikipedia-index.json http://localhost:8081/druid/indexer/v1/task`

Ingestion from Source Database 1/4

- Batch ingestion using *parallel_index* task type
- Sample OLTP database - AdventureWorks
 - Official Microsoft sample database
 - Several schemas: purchasing, sales, person, production
 - Sample DW [diagrams](#)
- Source database originally stored in *postgre* container
 - Run *load-data.sh* in order to prepare source relational database
- Ingestion task uses JDBC to retrieve data from database
 - *inputSource* type - sql
 - *connectorConfig* with connection string

Ingestion from Source Database 2/4

- Source tables from *production* schema
 - products, (sub)categories, work orders, work order routings
- Dimensions and metrics
 - Timestamp - e.g. start date of work order
 - Dimensions - product ID, product category
 - Metrics - ordered quantity, scraped quantity
 - Ingestion specs examples with and without metrics
- Granularity
 - Segment granularity in months
 - Query granularity in days
- Running ingestion task
 - `curl -X 'POST' -H 'Content-Type:application/json' -d @spec/spec-wo.json 'http://localhost:8888/druid/indexer/v1/task'`

Ingestion from Source Database 3/4

- Ingestion specs for

- Simple work order

```
SELECT startDate, workOrderId, productId, orderQty
FROM production.workorder
```

- Extended work order with product categories

```
SELECT workorderid, startdate, orderqty, w.productid, p.name as productname,
       psc.name as subcategory, pc.name as category
FROM production.workorder w
```

```
    JOIN production.product p ON w.productid = p.productid
```

```
    JOIN production.productssubcategory psc ON p.productssubcategoryid =
sc.productssubcategoryid
```

```
    JOIN production.productcategory pc ON psc.productcategoryid = pc.productcategoryid
```

- Work order routings

```
SELECT wo.workOrderId AS workOrderId, p.name AS productName, l.name AS locationName,
       wo.orderQty, operationSequence, plannedCost, actualCost, actualStartDate
```

```
FROM production.workOrderRouting wor
```

```
    JOIN production.location l ON wor.locationId = l.locationId
```

```
    JOIN production.product p ON wor.productId = p.productId
```

```
    JOIN production.workOrder wo ON wor.workOrderId = wo.workOrderId
```

Ingestion from Source Database 4/4

- Sample Druid SQL queries
- Work order
 - Work orders with largest quantity
 - Top 20 most ordered products
 - Products with multiple daily orders
 - Product with most *scrappedqty*
 - Weekly scrap rate
- Extended work order with categories
 - Top 10 most ordered bikes
 - Total ordered quantity per month
 - Most ordered subcategories
 - Subcategories with largest scrap rate
- Work order routing
 - Work orders with most routings
 - Locations sorted by routings
 - Painted products
 - Products with most resource hours

Querying and Analytics

- **Druid SQL**
 - Filtering, grouping, ordering
 - Joining datasources
 - SQL join
 - lookup
- **Native queries - JSON based**
 - Aggregation queries - TopN, GroupBy, Timeseries
 - Metadata queries
 - Other queries - Scan, Search

```
{
  "queryType": "topN",
  "dataSource": "sample_data",
  "dimension": "sample_dim",
  "threshold": 5,
  "metric": "count",
  "granularity": "all",
  "filter": {...},
  "aggregations": [...],
  "postAggregations": [...],
  "intervals": [...]
}
```


Ingestion - Individual exercises

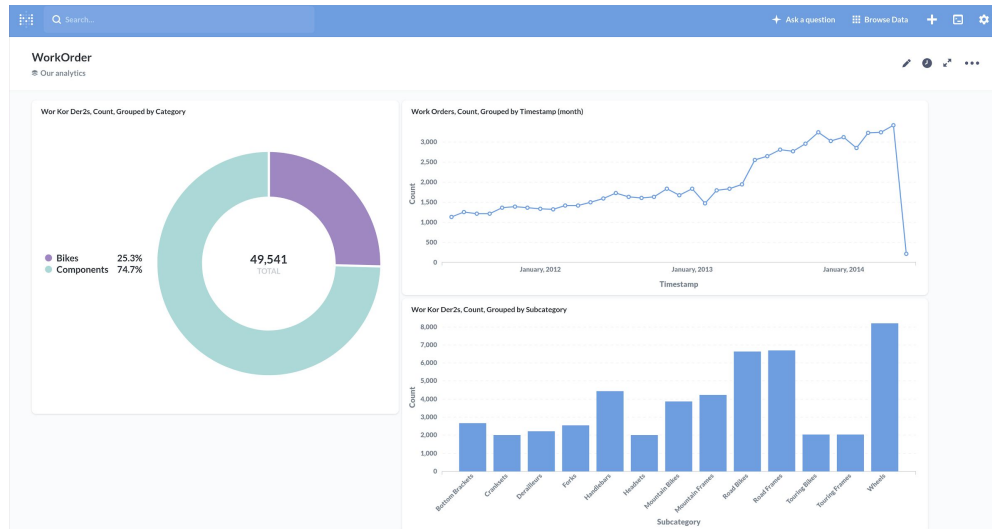
- Join multiple table from *sales* schema
 - SalesOrderDetail, SalesOrderHeader, Customer, SalesTerritory, SalesPerson, Product
- Define ingestion specs with metrics
 - Quantity, total amount, discount amount, tax amount, freight
- Submit ingestion task
- Monitor task execution
- Run some SQL queries on new datasource
 - Most popular products by territory
 - Customers by total amount
- Remove datasource

Apache Druid - Third Party Clients

- Druid's community developed several client libraries
 - Python, R, JavaScript, Clojure, Elixir, Ruby, SQL, PHP, Scala, Java, .NET, Rust
- Graphical UIs
 - Superset
 - Grafana
 - Pivot
 - Metabase
 - Metatron
- Other distributions, tools, and extensions

Metabase with Druid

- Web-based data analytics tool
- Connection to Druid broker
- Automatic retrieval of metadata
- User friendly
 - Asking questions using simple interface, complex custom queries, or native queries
 - Custom interactive dashboards with multiple graphs and diagrams based on questions
- Supports only Druid native queries



Metabase

Metabase with Druid

- Visit <http://localhost:3000>
- Create a new dashboard
- Ask question
- Add question to dashboard
- Place a graph/diagram on a dashboard
- Repeat the same steps for several questions

What's next?

- Streaming ingestion
 - Polling data from web (news, wikipedia, livescores)
 - Sending to Kafka
 - Druid ingests data stream from Kafka
 - Real-time analytics in Druid
- Batch ingestion
 - Hadoop data
- Complex ingestion
 - Transformations and filtering
- Druid Community events



That's all Folks!

Thank You!